

# Neural Networks & PyTorch

ICDSS - William Profit (wtp18)

# Overview

Mathematics of NNs

PyTorch framework

Data pipelines

Implementation

# Why?

Linear/Logistic regression can only model simple data distributions

Model more complicated functions

Desire to **loosely** mimic the brain

# Problem

Given a set of input-output pairs we want to find a function mapping input to output

E.g given:  $S = \{(x_1, y_1), (x_2, y_2) \dots (x_n, y_n)\}$

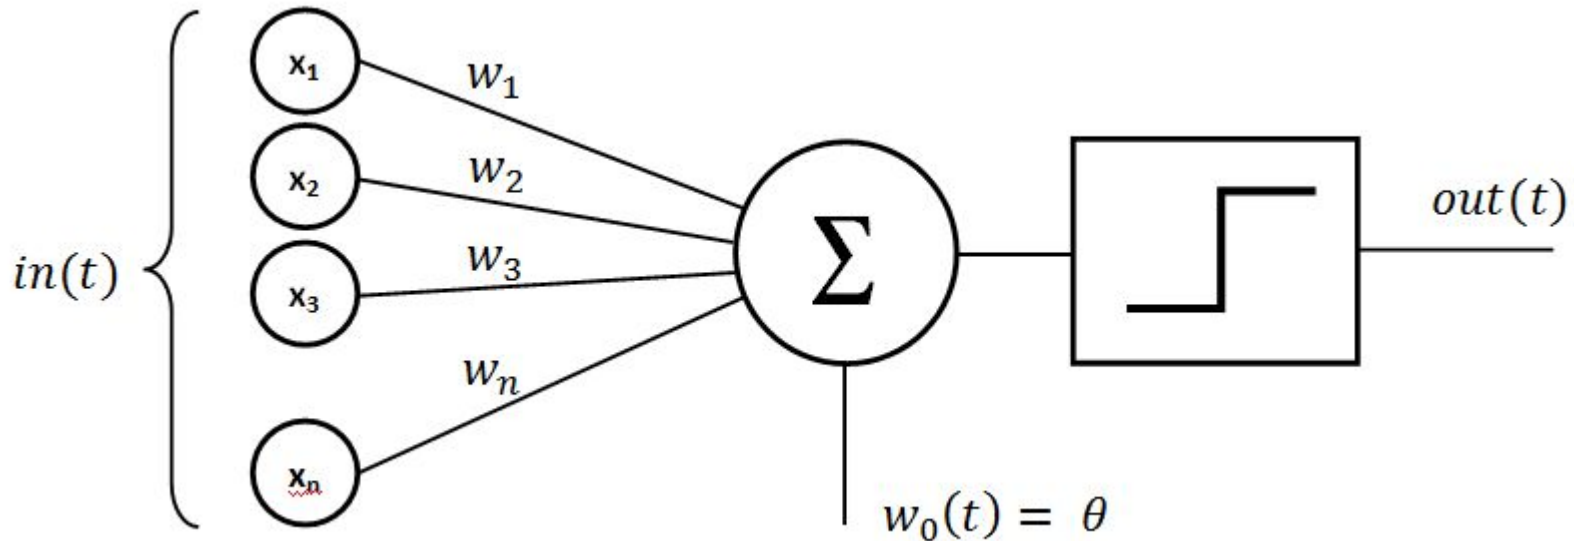
Find:  $f : \mathcal{X} \rightarrow \mathcal{Y}$

Example: map image (pixels) of scan to whether there is a tumor or not (binary classification)

# Perceptron

Example of a single neuron

Equivalent to logistic regression



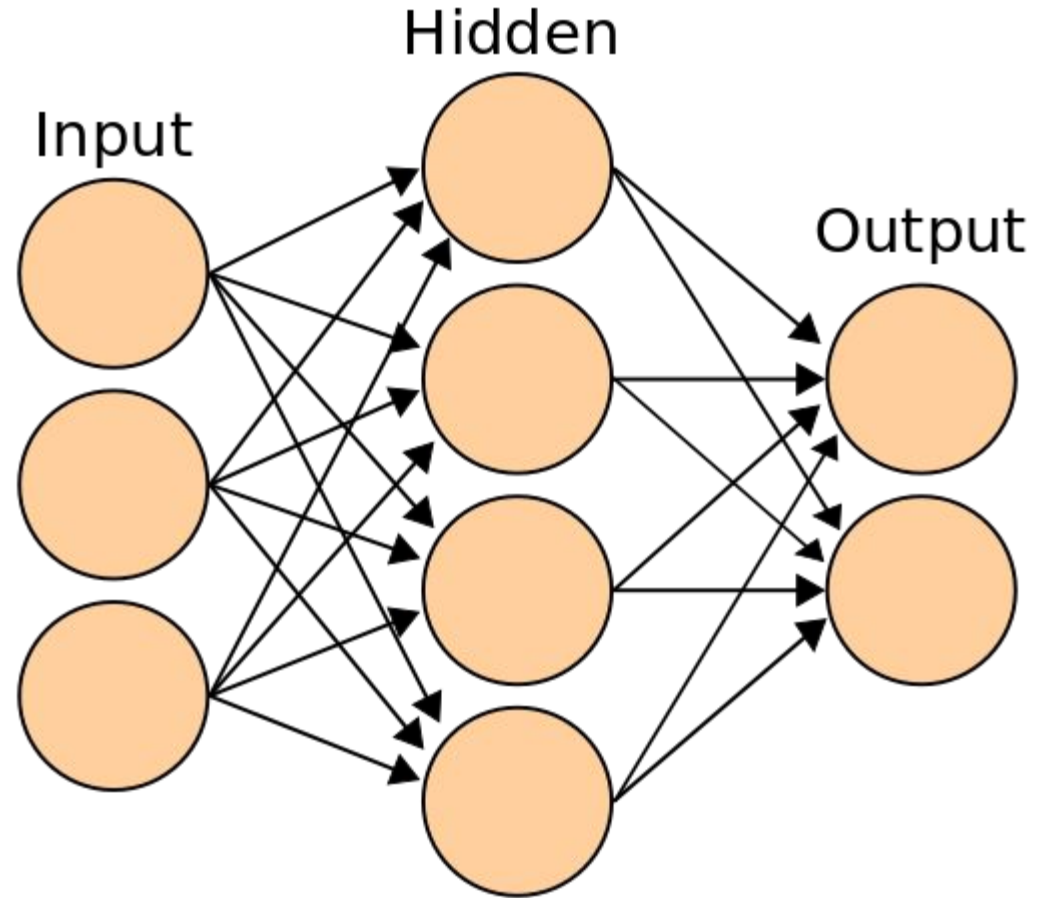
# Neural Network

Stack perceptrons into layers

Stack layers into network

Weights now represented by  
matrix

All parameters (weights)  
represented by  $\theta$



# Feed Forward

Every neuron computes its activation and forwards it to the next layer

Runs until the last layer, which gives us output

Predicted output is  $\hat{y} = Net(\theta, x)$

**=> Now need to train NN to predict correct value**

# Loss

Loss function  $\mathbf{L}(\boldsymbol{\theta})$  measures how  $\hat{y}$  differs from  $y$

Mean Squared Error for **regression**

$$\text{MSE} \triangleq \frac{1}{n} \sum_{i=1}^n (\hat{y}_i - y_i)^2$$

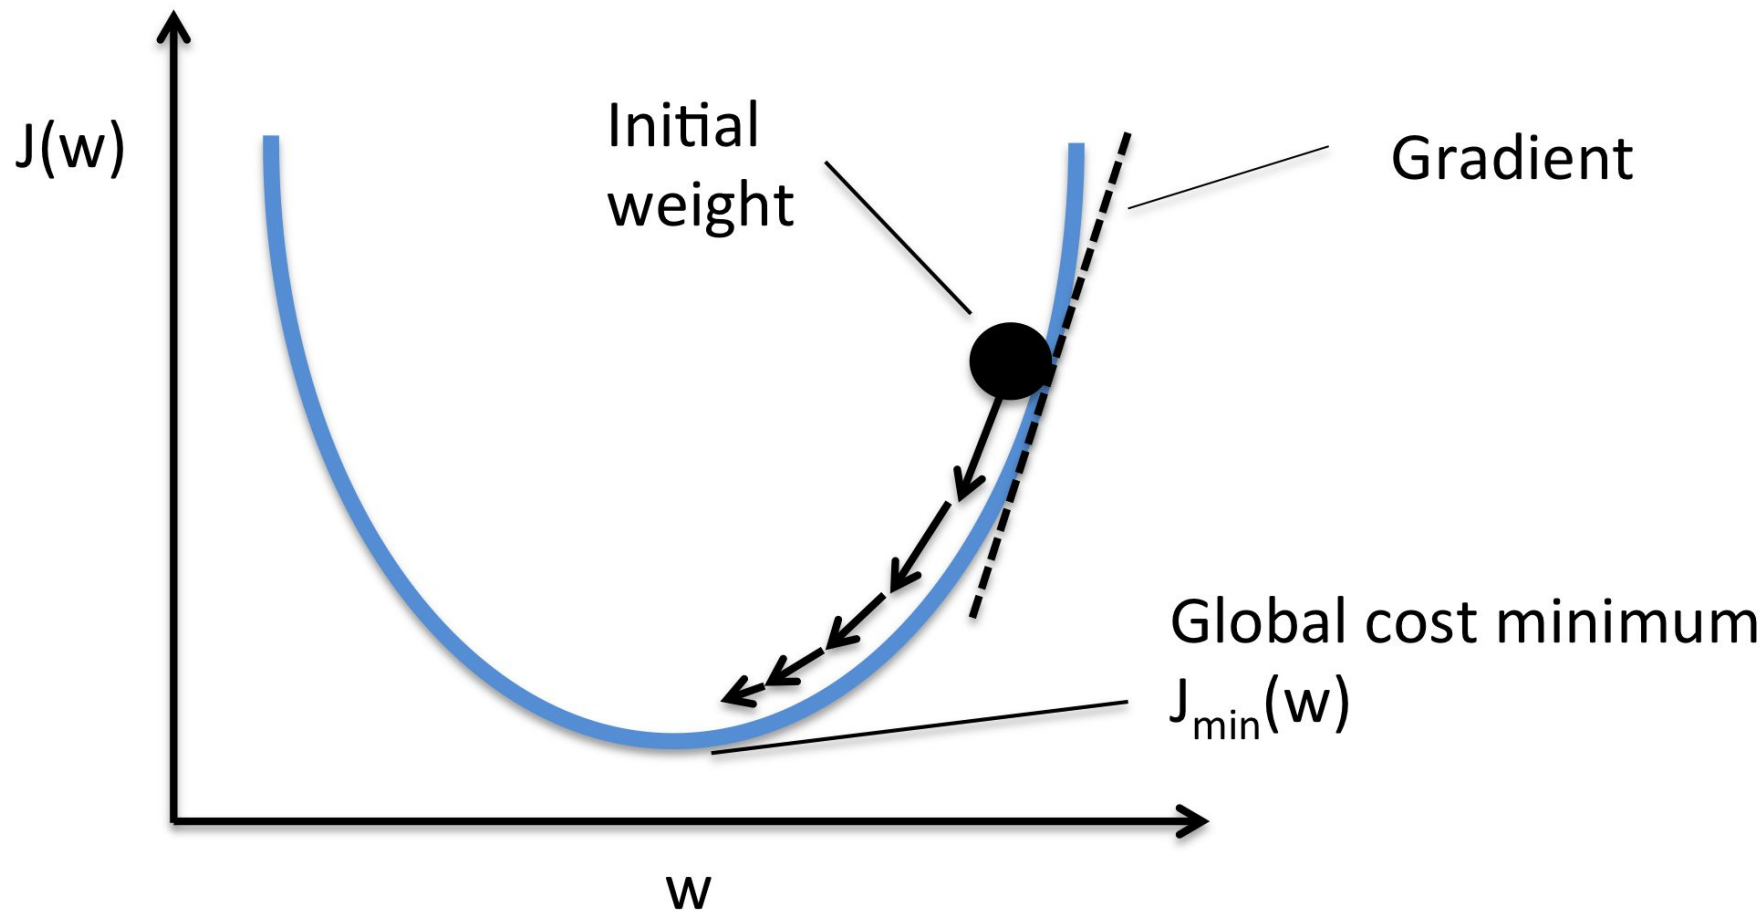
Cross Entropy for **binary classification**

$$\text{CE} \triangleq -\frac{1}{n} \sum_{i=1}^n y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)$$

Categorical Cross Entropy for **classification**

$$\text{CCE} \triangleq -\frac{1}{n} \sum_{i=1}^n \sum_{j=1}^m y_{ij} \log(\hat{y}_{ij})$$





# Backpropagation

We have now computed  $L(\theta)$  using the predicted  $\hat{\mathbf{y}}$  and the ground truth  $\mathbf{y}$

Now for every parameter, compute its gradient  $\frac{\partial L(\theta)}{\partial \theta}$

Computed using chain rule from last layer to first layer (i.e. propagating gradients back -> backpropagation)

=> For every parameter, we measure how it influences the loss.

Goal is now to minimise that loss using **Gradient Descent**

# Gradient Descent

For each parameter of  $\theta$ :

- Compute its gradient w.r.t. the loss (backprop)
- Update its value using learning rate alpha
- Repeat

$$\textit{repeat}\left\{\begin{array}{l}\theta := \theta - \alpha \frac{\partial L(\theta)}{\partial \theta}\end{array}\right\}$$

# Full training process

Initialise NN randomly

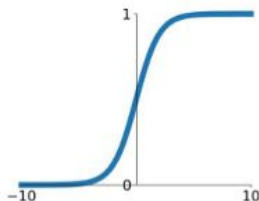
Repeat :

- Feed forward input
- Compute loss comparing output and ground truth
- Propagate gradients from loss back for each parameter in  $\theta$  (backprop)
- Update parameters with gradients and learning rate (gradient descent step)

# Activation Functions

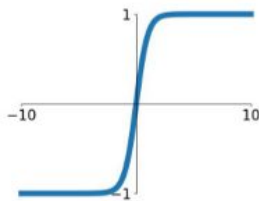
## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$



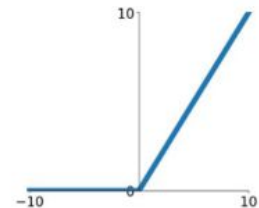
## tanh

$$\tanh(x)$$



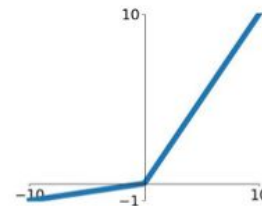
## ReLU

$$\max(0, x)$$



## Leaky ReLU

$$\max(0.1x, x)$$

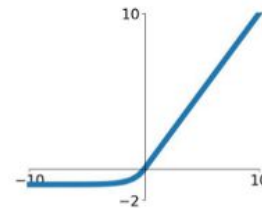


## Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

## ELU

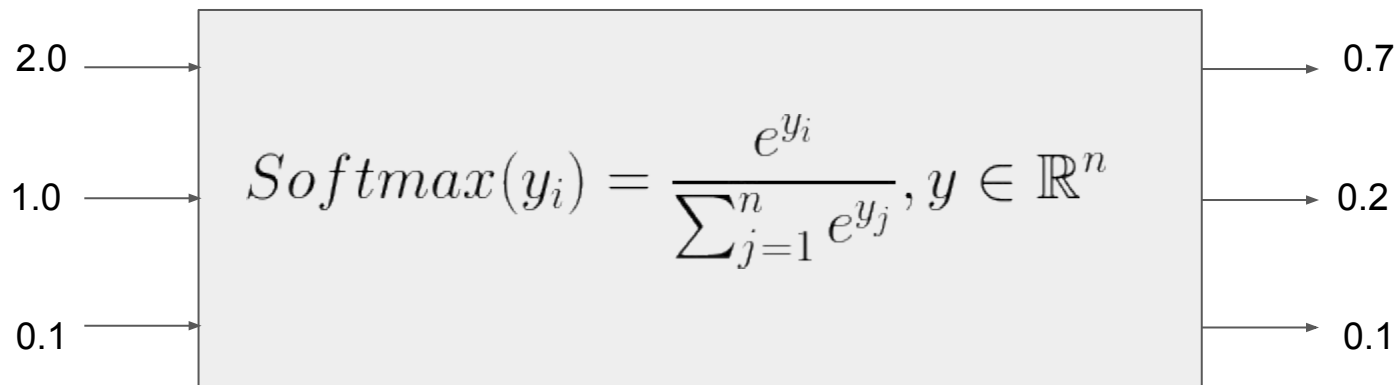
$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



# Softmax

Used to get probability distribution on output

All outputs sum to 1



# PyTorch

Tensors

Autograd

Neural Networks

Optimizers

Data Pipelines

# Tensors

```
import torch
import numpy as np

x = torch.tensor([[1, 2], [3, 4]], dtype=torch.float32)
y = torch.from_numpy(np.array([1, 2, 3, 4]))

x.to('cuda')

print(x.shape)
```



# Autograd

Used for automatic computation of gradients

```
x = torch.ones(2, 2, requires_grad=True)
y = x * x * 3
out = x.mean()

out.backward()
print(x.grad) #  $d(out)/dx$ 
```

# Neural Networks

```
import torch.nn as nn
import torch.nn.functional as F

class NeuralNet(nn.Module):
    def __init__(self):
        super(NeuralNet, self).__init__()
        self.fc1 = nn.Linear(784, 300)
        self.fc2 = nn.Linear(300, 10)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.softmax(self.fc2(x), dim=1)
```

# Optimising

```
import torch.optim as optim

loss_fn = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001)

optimizer.zero_grad()
output = net(input)
loss = loss_fn(output, ground_truth)
loss.backward()
optimizer.step()
```

# Datasets

Inherit ***Dataset***

Override **`__len__(self)`**

and **`__getitem__(self, index)`** methods

```
class MyDataset(Dataset):  
    def __init__(self, file, transform=None):  
        # Load data and do any preliminary operations  
        self.data = loadfile(file)  
        self.transform = transform  
  
    def __len__(self):  
        return len(self.data)  
  
    def __getitem__(self, index):  
        item = data[index]  
        item = self.transform(item)  
        return item
```

# Data Transforms

We can define transforms to be applied to our data before processing it

-> Used for data preprocessing pipeline

Examples: normalising, scaling, formatting, augmenting

Create class and define `__call__(self, sample)` method

```
class MyTransform():  
    def __init__(self, params):  
        some_initialisation_stuff(params)  
  
    def __call__(self, sample):  
        return transform_sample(sample)  
  
composed = transforms.Compose([Tranform1, Transform2])
```

# Data Loader

We now need to iterate through the dataset

Also want to shuffle the data (less biases)

Use batches for faster training

=> Use a data loader



```
dataset = MyDataset(  
    'path/to/data.csv',  
    transform.Compose([  
        Transform1,  
        Transform2  
    ])  
)  
  
dataloader = DataLoader(  
    dataset,  
    batch_size=32,  
    shuffle=True,  
    num_workers=4  
)
```

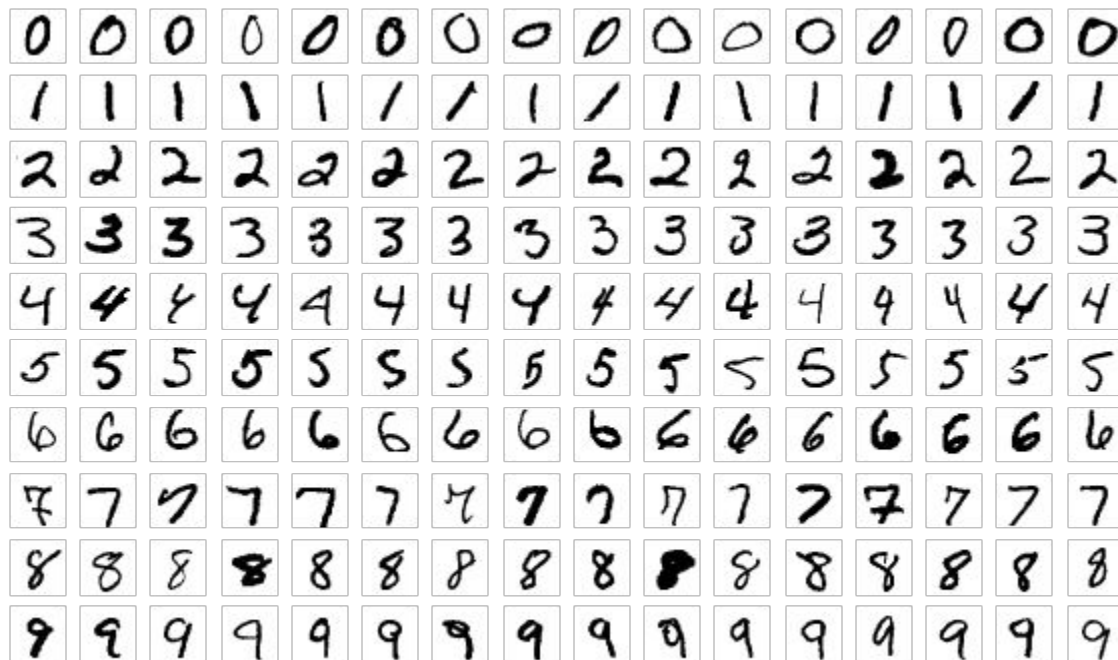
Iterate through data using a DataLoader

**data** is a tuple of the form [**inputs**, **labels**]

```
for i, data in enumerate(dataloader):  
    inputs, labels = data
```

# Demo

## MNIST Dataset - recognising handwritten digits



# Further reading

How backprop is computed (chain rule, computation graphs)

Regularisation, dropout

Different optimisers (adam, adadelta..)

Overfitting, underfitting

Gradient explosion/vanishing